

Learning General Policies for Planning through GPT Models

Nicholas Rossetti, Massimiliano Tummolo, Alfonso Emilio Gerevini,
Luca Putelli, Ivan Serina, Mattia Chiari, Matteo Olivato



UNIVERSITÀ
DEGLI STUDI
DI BRESCIA



LLMs still can't plan

- Models pre-trained on language datasets **cannot plan** using prompting techniques [Arora et al 2023, Valmeekam et al 2022, 2023]
- Via fine-tuning, Plansformer plans for various domains but **with a small number of objects** [Pallagani et al 2023a, 2023b]
- Does the problem lie on the **Transformer architecture** or on **the language pre-training dataset**?

Generative Pre-Trained Transformer (GPT)



- Once trained, given a sequence of initial words, GPT **predicts the next word** (e.g. names, verbs or adjectives)
- The predicted word **is added to the input** and GPT **repeats** the full process auto-regressively, obtaining the next word
- To be processed by GPT a sentence is divided into tokens (words) which are **embedded into a real number vector**

General Policies in Classical Planning



Given the current state and goal of the problem, a **general policy** is a function that provides the next action to apply:

$$\pi(s_i, Goal) = a \quad a \in A(s_i)$$

A policy π solves a set of classical planning instances for the same domain D if each of these instances $I = (O, Init, Goal)$ is **solved by the sequence of actions**:

$$\pi(s_0, Goal), \dots, \pi(s_n, Goal)$$

$$\text{where } s_0 = Init \text{ and } Goal \subseteq s_{n+1}$$

General Policies in Classical Planning



Given the current state and goal of the problem, a **general policy** is a function that provides the next action to apply:

$$\pi(Init, G, P) = a \quad a \in A(s) \quad s = \text{state produced by } P$$

A policy π solves a set of classical planning instances for the same domain D if each of these instances $I = (O, Init, Goal)$ is **solved by the sequence of actions**:

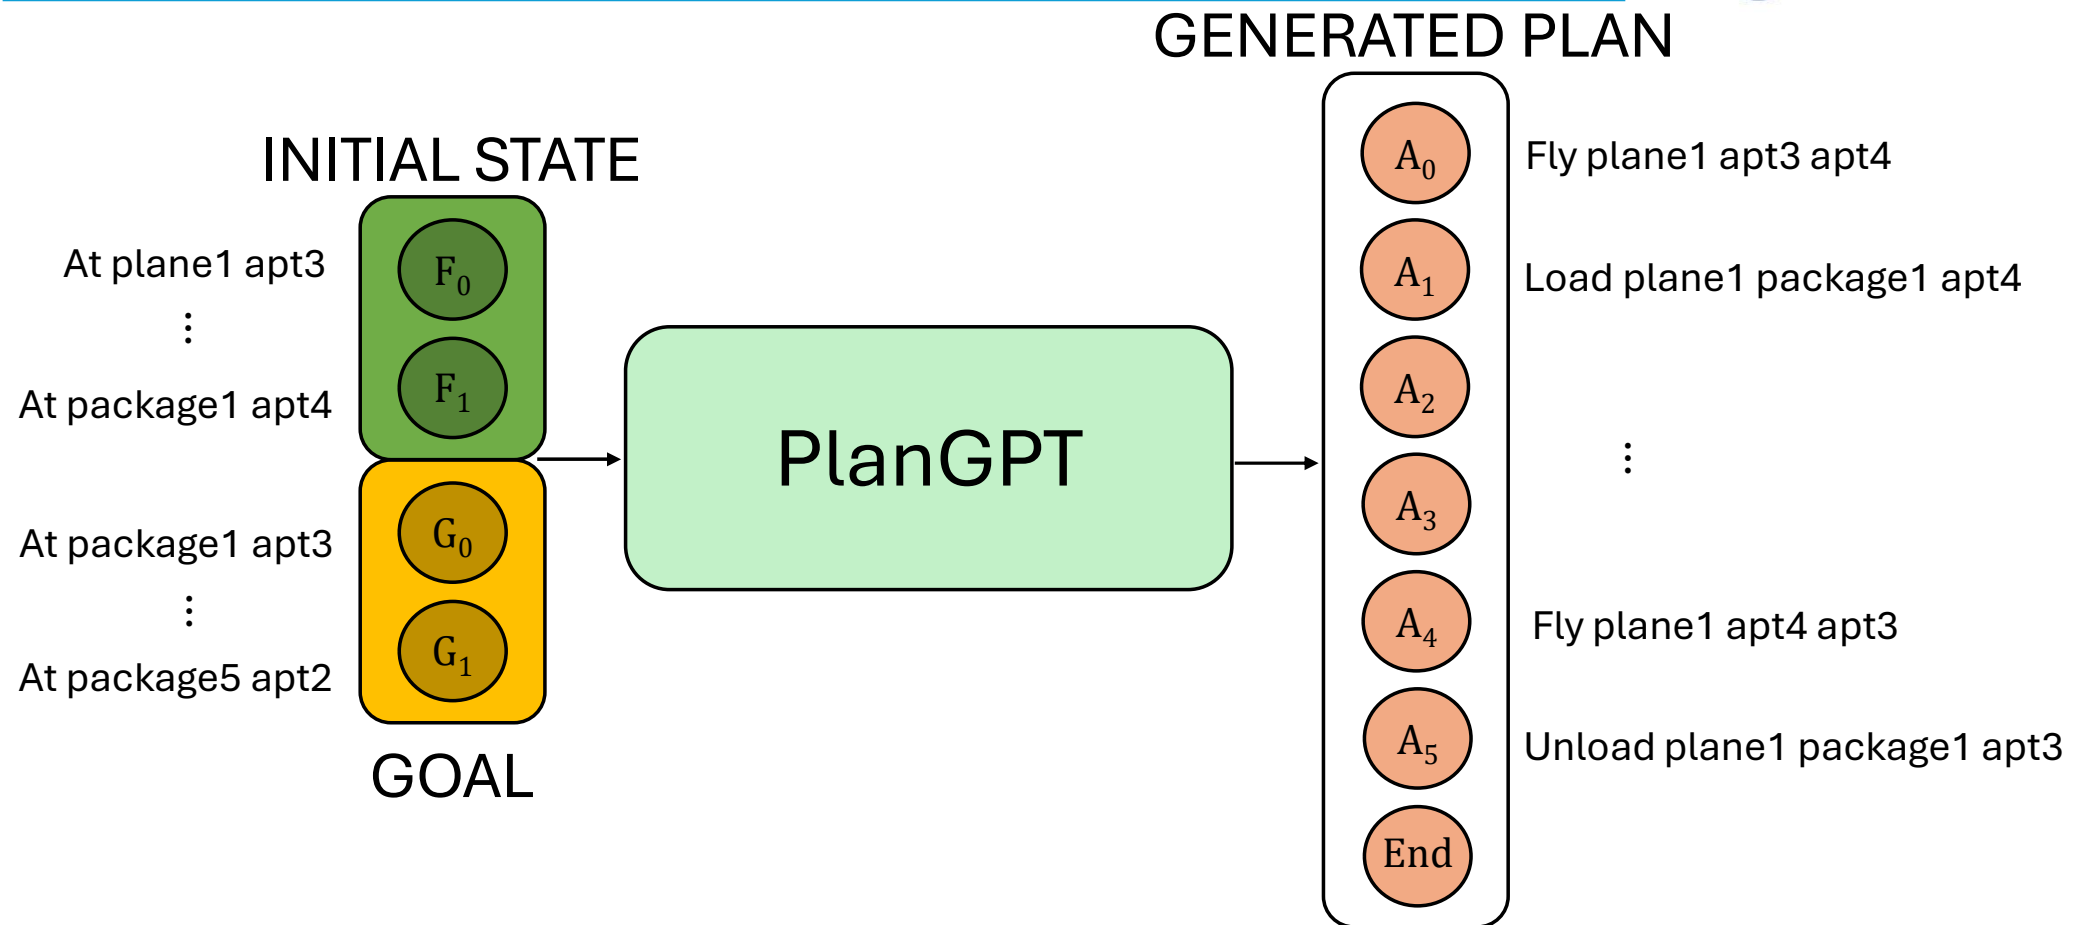
$$\pi(s_0, Goal), \dots, \pi(s_n, Goal)$$

$$\text{where } s_0 = Init \text{ and } Goal \subseteq s_{n+1}$$

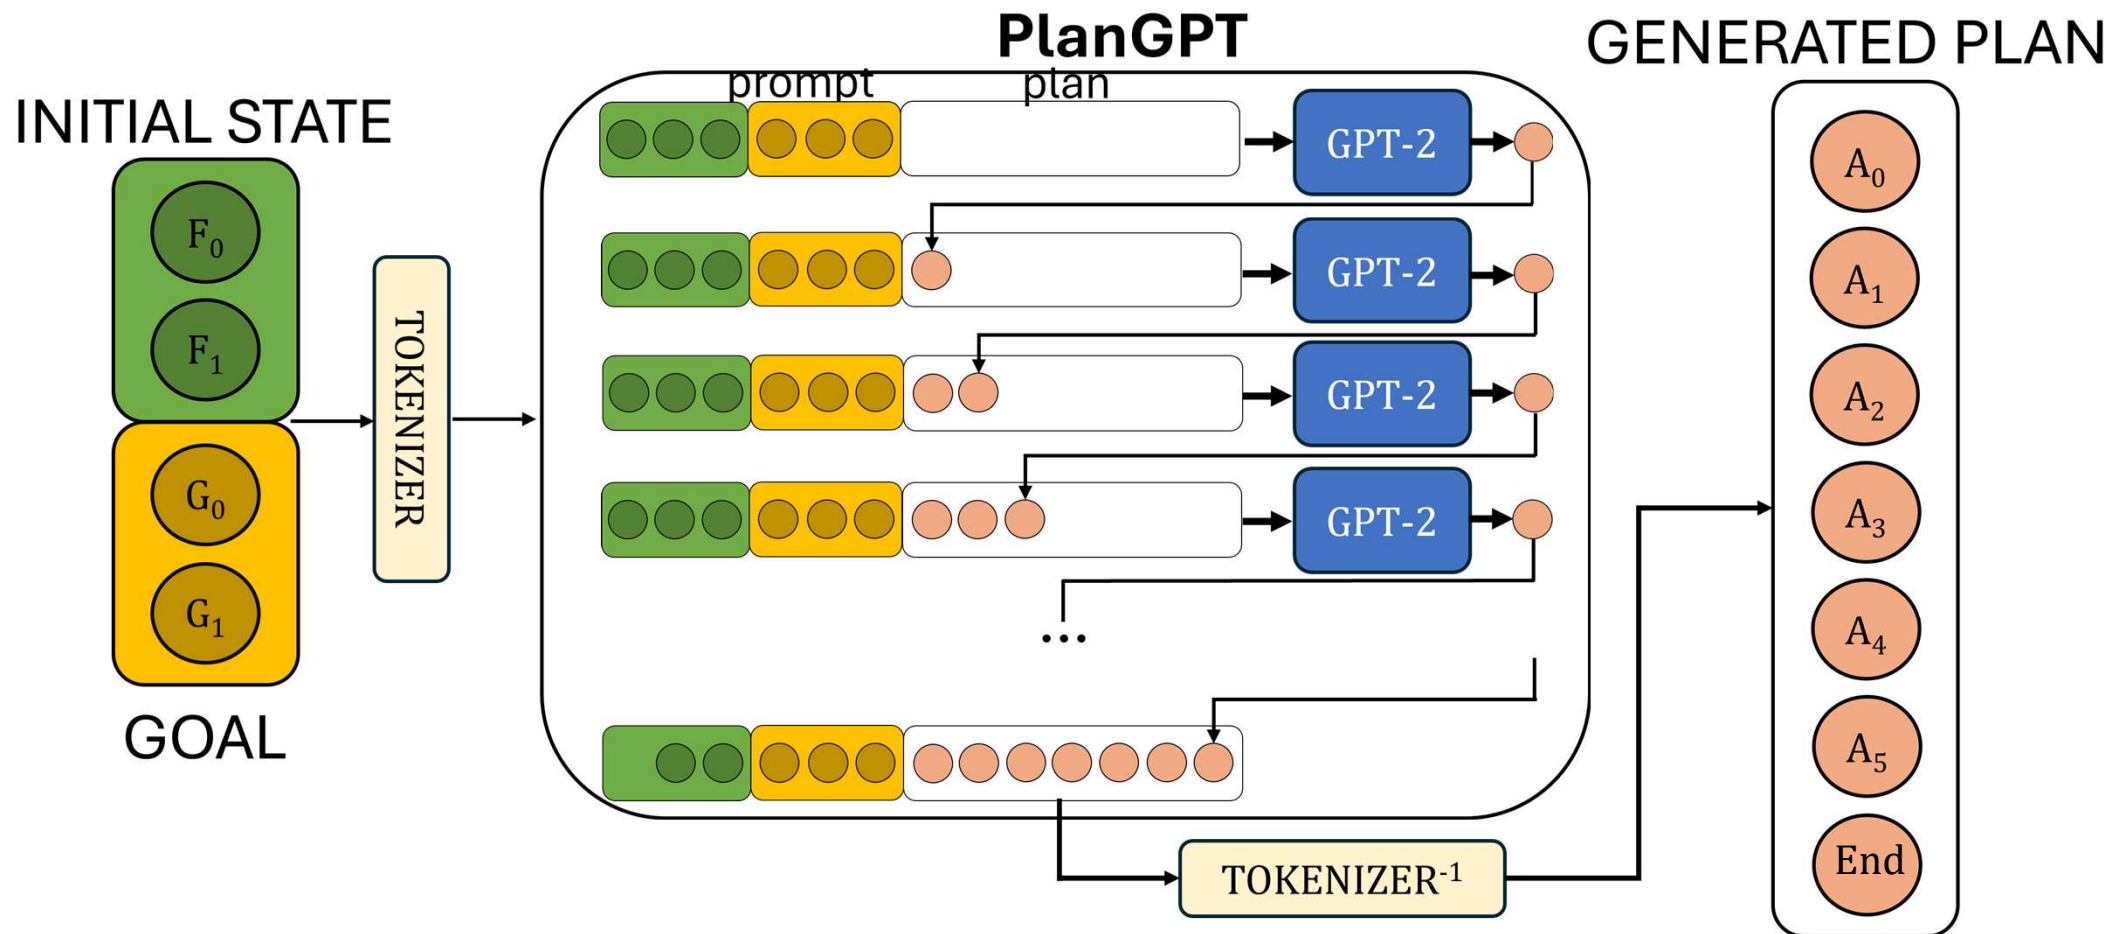
Can LLMs plan, if trained?

- We trained a GPT-2 model **from scratch** to obtain a **general policy**, using a dataset of **solved planning problems**
- We use the **initial state** and the **goal** of the problem as a prompt to GPT
- Given this prompt, we train GPT (**PlanGPT**) to **generate a valid plan**

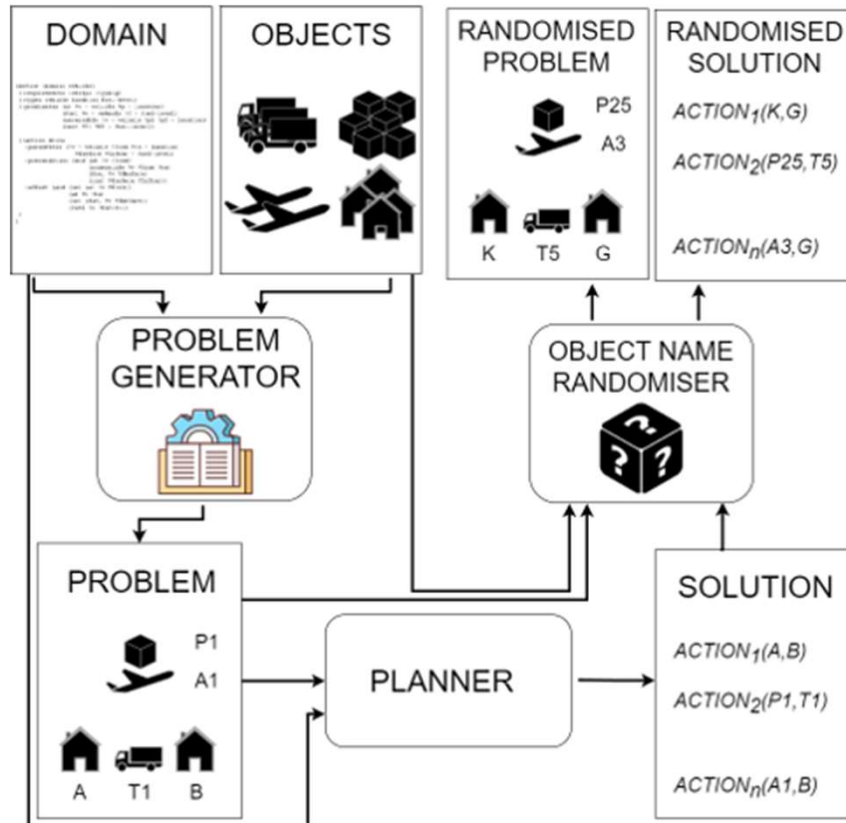
GPT for Planning (General Policies)



GPT for Planning (General Policies)



Dataset Generation



- We generated problems using a **problem generator**
- We solved these problems using **LPG** [Gerevini et al 2002]
- We randomize the objects' names

truck1 → truck4
city1 → city7

Training PlanGPT – the loss function

- The tokenizer splits each predicate and each action in their components: **predicate/action name and its objects**

at truck4 pos5	→	at - truck4 - pos5
drive truck4 pos5 pos7	→	drive - truck4 - pos5 – pos7

- We train PlanGPT by generating the next token and comparing it with the token in the solution plan using the **Cross Entropy loss**

$$L = - \sum_{t=1}^{|V|} \log(y'_t) * y_t$$

Cross Entropy Loss - Problem

- The use of the Cross Entropy Loss forces the learned model **to mimic the example label** (*a specific plan*)
- **Problem:** Valid plans different from the label plan are **considered incorrect** in the loss function!

Coverage Early Stopping

- **Solution:** We can mitigate this problem by *evaluating the model capability of generating valid plans during training*: let's include a **plan validator**!
- We evaluate our model based on the **coverage obtained on the validation set** instead of the validation loss using the **early stopping** technique.

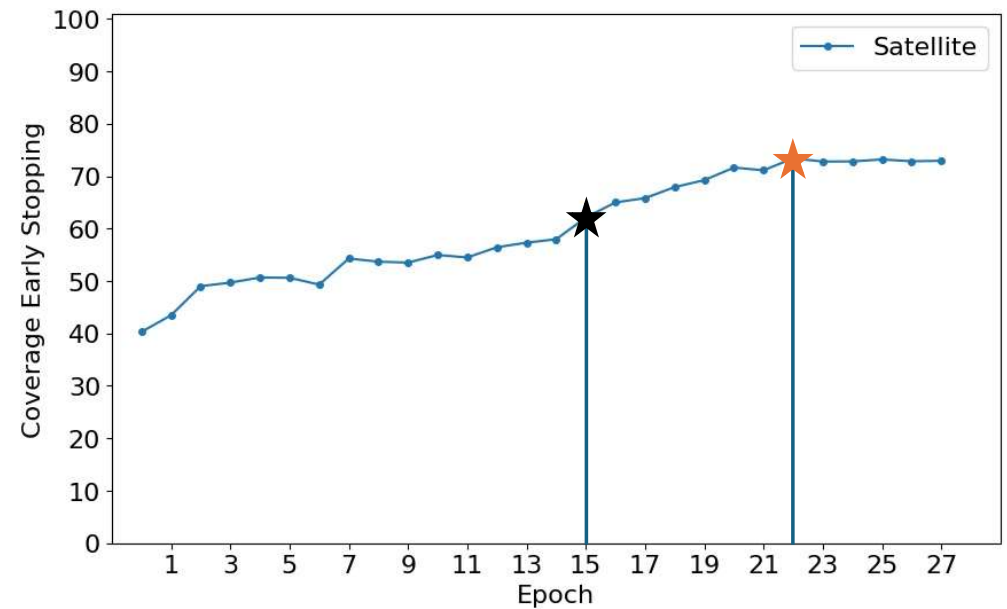

Experimental Setup

- We generate **70k problems**, divided in 63k for training set (obtaining more than **200k plans**), **1k** for validation set and **6k** for test set (**Tset**).
- We trained PlanGPT on **8 classical planning domains**: Blocksworld, Depots, Driverlog, Floortile, Logistics, Satellite, Visitall and Zenotravel
- We validate PlanGPT using classical planning metrics: **Coverage**, **IPC-Q** and **IPC-A**

Results – Coverage Early Stopping

Domain	Cross Entropy Loss Optimization	Coverage Early Stopping Optimization
Blocksworld	100.0	100.0
Depots	90.3	94.5
Driverlog	94.7	96.5
Floortile	98.2	99.6
Logistics	76.3	77.3
Satellite	81.3	90.1
Visitall	99.9	100.0
Zenotravel	94.7	94.7

Results – Coverage with Early Stopping



Comparison with Relational GNNs

- We compare with general policies obtained by Relational GNNs (Ståhlberg et al 2022a, 2022b):

Tset	Coverage		IPC-A		IPC-Q	
Domain	PlanGPT	R-GNN	PlanGPT	R-GNN	PlanGPT	R-GNN
Blocks	100.0	81.4	1763.1	1093.7	1847.1	1459.0
Logistics	77.3	21.6	4752.2	791.7	5125.1	772.1
Visitall	100.0	96.0	5754.5	3176.4	6046.4	6002.0

Conclusions

- We propose a method to learn a general policy based on GPT, called **PlanGPT**
- Our training procedure exploits an early stopping technique that we designed to increase the coverage
- For several domains, **PlanGPT** solves most of the problems in test sets, showing **competitive performance** w.r.t. SotA
- PlanGPT is limited to its **vocabulary** and **context window** and fails to generalize to problems **with more objects** compared to those in training

Future work

- We are exploring the use of PlanGPT has a heuristic: **providing a “good” plan seed to a plan-repair system (LPG)**
- We are investigating the integration of PlanGPT with **validator-driven RL techniques**, as in an RLHF setting
- We are investigating the integration of a **plan validator in the generation strategy** of PlanGPT

Thank you
Any questions?

Bibliography

- Arora, D.; and Kambhampati, S. 2023. Learning and Leveraging Verifiers to Improve Planning Capabilities of Pretrained Language Models. CoRR, abs/2305.17077
- Gerevini, A.; and Serina, I. 2002. LPG: A Planner Based on Local Search for Planning Graphs with Action Costs. In AIPS, 13–22. AAAI Press.
- Pallagani, V.; Muppasani, B.; Murugesan, K.; Rossi, F.; Srivastava, B.; Horesh, L.; Fabiano, F.; and Loreggia, A. 2023a. Understanding the Capabilities of Large Language Models for Automated Planning. CoRR, abs/2305.16151.
- Pallagani, V.; Muppasani, B.; Srivastava, B.; Rossi, F.; Horesh, L.; Murugesan, K.; Loreggia, A.; Fabiano, F.; Joseph, R.; and Kethepalli, Y. 2023b. Plansformer Tool: Demonstrating Generation of Symbolic Plans Using Transformers. In IJCAI, 7158–7162. IJCAI Org.
- Ståhlberg, S.; Bonet, B.; and Geffner, H. 2022a. Learning General Optimal Policies with Graph Neural Networks: Expressive Power, Transparency, and Limits. In ICAPS, 629–637. AAAI Press.
- Ståhlberg, S.; Bonet, B.; and Geffner, H. 2022b. Learning Generalized Policies without Supervision Using GNNs. In KR, 474–483. IJCAI Org.
- Valmeekam, K.; Hernandez, A. O.; Sreedharan, S.; and Kambhampati, S. 2022. Large Language Models Still Can't Plan (A Benchmark for LLMs on Planning and Reasoning about Change). CoRR, abs/2206.10498.
- Valmeekam, K.; Sreedharan, S.; Marquez, M.; Hernandez, A. O.; and Kambhampati, S. 2023. On the Planning Abilities of Large Language Models (A Critical Investigation with a Proposed Benchmark). CoRR, abs/2302.06706.